Solving the TTC Model Execution Case with FunnyQT

Tassilo Horn

Institute for Software Technology, University Koblenz-Landau, Germany

horn@uni-koblenz.de

This paper describes the FunnyQT solution to the TTC 2015 Model Execution transformation case. The solution solves the third variant of the case, i.e., it considers and implements the execution semantics of the complete UML Activity Diagram language.

FunnyQT is a model querying and model transformation library for the functional Lisp-dialect Clojure providing a comprehensive and efficient querying and transformation API, many parts of which are provided as task-oriented embedded DSLs.

1 Introduction

This paper describes the FunnyQT[1] [2, 3] solution of the TTC 2015 Model Execution Case [4]. It implements the third variant of the case description, i.e., it implements the execution semantics of the complete UML Activity Diagram language. The solution project is available on Github[2], and it is set up for easy reproduction on a SHARE image[3].

FunnyQT is a model querying and transformation library for the functional Lisp dialect Clojure[4]. Queries and transformations are plain Clojure programs using the features provided by the FunnyQT API.

As a Lisp, Clojure provides strong metaprogramming capabilities that are exploited by FunnyQT in order to define several *embedded domain-specific languages* (DSL, [1]) for different querying and transformation tasks.

FunnyQT is designed with extensibility in mind. By default, it supports EMF [5] models and JGraLab[5] TGraph models. Support for other modeling frameworks can be added without having to touch FunnyQT's internals.

The FunnyQT API is structured into the following namespaces, each namespace providing constructs supporting concrete querying and transformation use-cases:

funnyqt.emf EMF-specific model management API

funnyqt.tg JGraLab/TGraph-specific model management API

funnyqt.generic Protocol-based, generic model management API

funnyqt.query Generic querying constructs such as quantified expressions or regular path expressions

funnyqt.polyfns Constructs for defining polymorphic functions dispatching on metamodel types

funnyqt.pmatch Pattern matching constructs

funnyqt.relational Constructs for logic-based, relational model querying inspired by Prolog

funnyqt.in-place In-place transformation rule definition constructs

---

[1] http://funnyqt.org
[2] https://github.com/tsdh/ttc15-model-execution-funnyqt
[3] FIXME:addSHAREURL
[4] http://clojure.org
[5] http://jgralab.github.io

funnyqt.model2model Out-place transformation definition constructs similar to ATL or QVT Operational Mappings

funnyqt.extensional Transformation API similar to GReTL

funnyqt.bidi Constructs for defining bidirectional transformations similar to QVT Relations

funnyqt.coevo Constructs for transformations that evolve a metamodel and a conforming model simultaneously at runtime

funnyqt.visualization Model visualization

funnyqt.xmltg Constructs for querying and modifying XML files as models conforming to a DOM-like metamodel

For solving the model execution case, the *funnyqt.emf*, *funnyqt.query*, and *funnyqt.polyfns* namespaces have been used.

2   Solution Description

The explanations in the case description about the operational semantics on UML Activity Diagrams suggest an algorithmic solution to the transformation case. The FunnyQT solution tries to be almost a literal translation of the case description to Clojure code.

The first line of the solution calls the `generate-ecore-model-functions` FunnyQT macro.

```
1 (generate-ecore-model-functions "activitydiagram.ecore" ttc15-model-execution-funnyqt.ad a)
```

As its name suggests, it generates a metamodel-specific API. This API is generated into the namespace `ttc15-model-execution-funnyqt.ad`, and the namespace alias `a` is used to refer to that namespace from the current one.

The generated API consists of element creation functions, lazy element sequence functions, attribute access functions, and reference access functions. For example, (`a/create-ControlToken!` ad) creates a new control token and adds it to the activity diagram model `ad`, (`a/eall-Inputs` ad) returns the lazy sequence of input elements in `ad`, (`a/running?` n) and (`a/set-running!` n true) query and set the node `n`'s running attribute, and (`a/->locals` a), (`a/->set-locals!` a ls), and (`a/->add-locals!` a l) query, set, and add to the locals reference of the activity `a`.

Metamodel-specific APIs created by the `generate-ecore-model-functions` macro are independent of the underlying implementation kind. The FunnyQT solution uses a dynamic instance model but it would work equivalently if it was run with an EMF model using classes generated for the metamodel's generator model.

Instead of using the generated API, the solution could also use the generic EMF API in which case the lazy sequence of a activity diagram models inputs would be retrieved using (`eallcontents` ad 'Input). But the generated API is slightly more concise and readable.

In the following, the solution is presented in a top-down manner similar to how the case description defines the operational semantics of activity diagrams. Listing 1 on the facing page shows the function `execute-activity-diagram` which contains the transformation's main loop.

The function queries the single activity in the diagram, creates a new trace, and assigns that to the activity. Then, the activity's variables are initialized and its nodes are set running.

Then, a `loop-recur` iteration[6] performs the actual execution of the activity. The variable `ens` is bound to the enabled nodes. Initially, those are the activity's initial nodes from which there can only be one. As

---

[6] `loop` is not a loop in the sense of Java's `for` or `while` but a local tail-recursion construct. The `loop` declares variables with their initial bindings, and in the `loop`'s body `recur` forms may recurse back to the beginning of the `loop` providing new bindings for the `loop`'s variables.

```
2 (defn execute-activity-diagram [ad]
3   (let [activity (the (a/eall-Activities ad))
4         trace (a/create-Trace! nil)]
5     (a/->set-trace! activity trace)
6     (init-variables activity (first (a/eall-Inputs ad)))
7     (mapc #(a/set-running! % true) (a/->nodes activity))
8     (loop [ens (filter a/isa-InitialNode? (a/->nodes activity))]
9       (when (seq ens)
10        (doseq+ [node ens]
11          (exec-node node)
12          (a/->add-executedNodes! trace node))
13        (recur (enabled-nodes activity))))
14    trace))
```

Listing 1: The main loop of executing activity diagrams

long as there are enabled nodes left, each node gets executed one after the other and then added to the trace. Thereafter, the loop is restarted with the nodes that are enabled at that point in time. Eventually, there won't be any enabled nodes left, and then the function returns the trace.

So the first step in the execution of an activity is the initialization of its local and input variables. The corresponding function `init-variables` is shown in listing 2. For locals, their current value is set to their initial value if there is one defined. For input variables, their current value is set to the value of the input's corresponding input value element.

```
15 (defn init-variables [activity input]
16   (doseq+ [lv (a/->locals activity)]
17     (when-let [init-value (a/->initialValue lv)]
18       (a/->set-currentValue! lv init-value)))
19   (doseq+ [iv (and input (a/->inputValues input))]
20     (when-let [val (a/->value iv)]
21       (a/->set-currentValue! (a/->variable iv) val))))
```

Listing 2: Initialization of variables

After initializing the variables, the main function sets the activity's nodes running, and the main loop starts with the activity's initial node which will be executed.

For different kinds of activity nodes, different execution semantics have to be encoded. This is exactly the use-case of FunnyQT's polymorphic functions (polyfn). A polymorphic function is declared once, and then different implementations for instances of different metamodel types can be defined. When the polyfn is called, a polymorphic dispatch based on the polyfn's first argument's metamodel type is performed to pick out the right implementation[7].

Listing 3 on the following page shows the declaration of the polyfn `exec-node` and its implementation for initial nodes. The declaration only defines the name of the polyfn and the number of its arguments (just one, here). The implementation for initial nodes simply offers one new control token to the initial node's outgoing control flow edge.[8]

Listing 4 on the next page shows the `exec-node` implementations for join, merge, and decision nodes. Join and Merge nodes simply consume their input offers and pass the tokens they have been offered on all outgoing control flows. Decision nodes act similar but offer their input tokens only on the outgoing control flow whose guard variable's current value is true[9].

---

[7]Polyfns support multiple inheritance. In case of an ambiguity during dispatch, e.g., two or more inherited implementations are applicable, an error is signaled.

[8]The FunnyQT function `the` is similar to Clojure's `first` except that it signals an error if the given collection contains zero or more than one element. Thus, it makes the assumption that there must be only one outgoing control flow explicit.

[9]`(the predicate collection)` returns the single element of the collection for which the predicate returns true. If

```
22 (declare-polyfn exec-node [node])

23 (defn offer-one-ctrl-token [node]
24   (let [ctrl-t (a/create-ControlToken! nil)
25         out-cf (the (a/->outgoing node))
26         offer  (a/create-Offer! nil {:offeredTokens [ctrl-t]})]
27     (a/->add-heldTokens! node ctrl-t)
28     (a/->add-offers! out-cf offer)))

29 (defpolyfn exec-node InitialNode [i]
30   (offer-one-ctrl-token i))
```

Listing 3: Declaration of polyfn `exec-node` and its implementation for initial nodes

```
31 (defn pass-tokens
32   ([n] (pass-tokens n nil))
33   ([n out-cf]
34    (let [in-toks (consume-offers n)]
35      (a/->set-heldTokens! n in-toks)
36      (doseq+ [out-cf (if out-cf [out-cf] (a/->outgoing n))]
37        (a/->add-offers!
38          out-cf (a/create-Offer!
39                   nil {:offeredTokens in-toks}))))))

40 (defpolyfn exec-node JoinNode [jn]
41   (pass-tokens jn))

42 (defpolyfn exec-node MergeNode [mn]
43   (pass-tokens mn))

44 (defpolyfn exec-node DecisionNode [dn]
45   (pass-tokens dn (the #(-> % a/->guard a/->currentValue a/value)
46                        (a/->outgoing dn))))
```

Listing 4: `exec-node` impls for join, merge, and decision nodes

So how does a node consume offers? This is defined by the `consume-offers` function shown in listing 5. First, the offers and the tokens offered by them are calculated. Then, the offered tokens are divided into control and forked tokens. For control tokens, their holder is unset. For forked tokens, the corresponding base token's holder is unset. The forked tokens' remainingOffersCount is decremented. If it has become zero then, the forked token is removed from its holder. Lastly, the offers are deleted, and the incoming tokens are returned.

```
47 (defn consume-offers [node]
48   (let [offers    (mapcat a/->offers (a/->incoming node))
49         tokens    (mapcat a/->offeredTokens offers)
50         ctrl-toks (filter a/isa-ControlToken? tokens)
51         fork-toks (filter a/isa-ForkedToken? tokens)]
52     (doseq+ [ct ctrl-toks]
53       (a/->set-holder! ct nil))
54     (doseq+ [ft fork-toks]
55       (when-let [bt (a/->baseToken ft)]
56         (a/->set-holder! bt nil))
57       (a/set-remainingOffersCount! ft (dec (a/remainingOffersCount ft)))
58       (when (zero? (a/remainingOffersCount ft))
59         (a/->set-holder! ft nil)))
60     (mapc edelete! offers)
61     tokens))
```

Listing 5: Consuming offers

The remaining kinds of activity nodes are activity final nodes and opaque actions. Their `exec-node` implementations are printed in listing 6 on the next page. An activity final node simply consumes all

---

there is no or more elements satisfying the predicate, an error is signaled.

offers and then sets the running attribute of all nodes contained by the executed activity to false. An opaque action also consumes all offers, then evaluates all its expressions in sequence using the `eval-exp` function, and finally offers one single control token on the outgoing control flow.

```
62 (defpolyfn exec-node ActivityFinalNode [afn]
63   (consume-offers afn)
64   (mapc #(a/set-running! % false)
65         (-> afn a/->activity a/->nodes)))

66 (defpolyfn exec-node OpaqueAction [oa]
67   (consume-offers oa)
68   (mapc eval-exp (a/->expressions oa))
69   (offer-one-ctrl-token oa))
```

Listing 6: `exec-node` impls for activity final nodes and opaque actions

How an expression is evaluated depends on (1) the type of the expression, and (2) on the value of the operator attribute. The expression's type is only important in order to separate unary from binary expressions, and the operator defines the semantics. Therefore, the `eval-exp` function shown in listing 7 has a special case for boolean unary expressions which negates the expression's current value using `not`. For all binary expressions, the map `op2fn` mapping from operator enum constants to Clojure functions having the semantics of that operator is used. The function determined by looking up the expression's operator is applied to both operands to compute the new value.

```
70 (def op2fn {(a/eenum-IntegerCalculationOperator-ADD)             +
71             (a/eenum-IntegerCalculationOperator-SUBRACT)         -
72             (a/eenum-IntegerComparisonOperator-SMALLER)          <
73             (a/eenum-IntegerComparisonOperator-SMALLER_EQUALS)   <=
74             (a/eenum-IntegerComparisonOperator-EQUALS)           =
75             (a/eenum-IntegerComparisonOperator-GREATER_EQUALS)   >=
76             (a/eenum-IntegerComparisonOperator-GREATER)          >
77             (a/eenum-BooleanBinaryOperator-AND)                  #(and %1 %2)
78             (a/eenum-BooleanBinaryOperator-OR)                   #(or  %1 %2)})

79 (defn eval-exp [exp]
80   (a/set-value! (-> exp a/->assignee a/->currentValue)
81                 (if (a/isa-BooleanUnaryExpression? exp)
82                   (not (-> exp a/->operand a/->currentValue a/value))
83                   ((op2fn (a/operator exp))
84                    (-> exp a/->operand1 a/->currentValue a/value)
85                    (-> exp a/->operand2 a/->currentValue a/value)))))
```

Listing 7: Evaluation of expressions

After executing all enabled nodes, the transformation's main function `execute-activity-diagram` from listing 1 on page 3 recomputes the enabled nodes and resumes the execution. The enabled nodes are computed by the `enabled-nodes` function shown in listing 8 on the following page. The enabled nodes are those nodes of a given activity which are set running, are no initial nodes[10], and receive an offer on each incoming control flow, or, in the case of a merge node, on one incoming control flow.

These 92 NCLOC of algorithmic FunnyQT/Clojure code implement the complete operational semantics of UML Activity Diagrams (with the exception of data flows which has not been demanded by the case description).

3   Evaluation

In this section, the FunnyQT solution is evaluated according to the criteria stated in the case description.

---

[10]Initial nodes have to be excluded because if they are set running, all of their (zero) incoming control flows have offers.

```
86 (defn enabled-nodes [activity]
87   (filter (fn [n]
88            (and (a/running? n)
89                 (not (a/isa-InitialNode? n))
90                 ((if (a/isa-MergeNode? n) exists? forall?)
91                  #(seq (a/->offers %)) (a/->incoming n))))
92           (a/->nodes activity)))
```

Listing 8: Computation of enabled nodes

Correctness.   The JUnit test suite of the reference Java solution has been translated to Clojure and its unit testing library. The very same assertions are tested, i.e., the execution order of nodes is tested and the final values of variables are checked. Additionally, it is checked that after an activity has been executed, there are no tokens leftover whose existence would hint at some bug in the implementation.

   All tests pass for all provided test models, so the solution is correct at least with respect to these models and the set of assertions tested by the unit test.

Understandability and conciseness.   With 92 lines of non-commented source code, the FunnyQT solution is quite concise.

   Of course, understandability is a very subjective measure measure. The solution should be evident for any Clojure programmer but even without prior Clojure knowledge, the solution shouldn't be hard to follow due to the usage of the metamodel-specific API. Another strong point is that all steps in the execution of an activity are encoded in one function each whose definition is almost a literal translation of the English description to FunnyQT/Clojure code.

Performance.   Table 1 shows the execution times of the FunnyQT solution for all provided test models. These times were measured on a normal 4-core laptop with 2.6 GHz and 2 GB of RAM dedicated to the JVM.

Model	Execution time
test1	1.4 ms
test2	0.8 ms
test3	3.1 ms
test4	4.3 ms
test5	0.7 ms
test6 (false)	6.5 ms
test6 (true)	10.8 ms
test-performance-variant-1	2093.3 ms
test-performance-variant-2	248.5 ms
test-performance-variant-3-1	118.2 ms
test-performance-variant-3-2	172.8 ms

Table 1: Execution times for the provided test models

   When compared with the reference Java solution, there are no notable differences for the smaller models. For the model *test-performance-variant-1*, the Java solution is slightly faster. However, for the model *test-performance-variant-2* the Java solution takes about 1600 ms, and for the *test-performance-variant-3-1* it takes about 8000 ms. So for these models, the FunnyQT solution is an order of magnitudes

faster than the Java solution which hints at some inefficiencies in how parallel branches are executed in the latter.

4 Conclusion

In this paper, the FunnyQT solution to the TTC 2015 Model Execution case has been discussed. The solution implements the full operational semantics of UML Activity Diagrams with the exception of object flows which haven't been considered in this case.

The specification is very concise. The complete implementation amounts to only 92 lines of code without counting comments and empty lines.

The specification is also well-understandable. The description of the semantics given in the case description have been translated almost literally to functions and one polymorphic function. These functions use a metamodel-specific API which FunnyQT has generated from the activity diagram metamodel.

The FunnyQT solution also performs very well. In general, its execution times are very similar to those of the Java reference solution, and for two of the performance test models, the FunnyQT solution is in fact several times faster.

Overall, FunnyQT seems to be very adequate for defining model interpreters. Especially its polymorphic function facility has been explicitly designed for these kinds of tasks.

References

[1] Martin Fowler (2010): *Domain-Specific Languages*. Addison-Wesley Professional.

[2] Tassilo Horn (2013): *Model Querying with FunnyQT - (Extended Abstract)*. In Keith Duddy & Gerti Kappel, editors: *ICMT, Lecture Notes in Computer Science* 7909, Springer, pp. 56–57.

[3] Tassilo Horn (2015): *Graph Pattern Matching as an Embedded Clojure DSL*. In: *International Conference on Graph Transformation - 8th International Conference, ICGT 2015, L'Aquila, Italy, July 2015*. To appear.

[4] Tanja Mayerhofer & Manuel Wimmer (2015): *The TTC 2015 Model Execution Case*. In: *Transformation Tool Contest 2015*.

[5] Dave Steinberg, Frank Budinsky, Marcelo Paternostro & Ed Merks (2008): *EMF: Eclipse Modeling Framework*, 2 edition. Addison-Wesley Professional.